

WHITE PAPER · GOVERNANCE PAPER

AI Safety and Professional Boundaries in Accounting

Where AI assistance ends and professional judgment begins, and how Kriton™ is designed to clarify, limit or escalate rather than answer definitively.

● AI Governance

● Audit

● Legal

FORMAT

PDF · Free download

READING TIME

About 16 minutes

EDITION

Q4 2026

Educational research only. This paper does not constitute accounting, tax, legal, audit or compliance advice.

ABOUT THIS PAPER

Read this first.

This paper is for the people who set the rules for AI use in professional accounting work. It treats safety as a design problem about boundaries, not a property of a model, and gives you a way to specify, test and monitor those boundaries. It is vendor-neutral, so you can apply it to any approach, including ours. It is not legal advice, and legal readers should take it as input to their own analysis.

WHO IT IS FOR	AI governance owners, audit and risk leaders, legal and compliance teams, and the professionals who will rely on AI assistance in accounting work.
THE QUESTION	Where does AI assistance end and professional judgment begin?
THE THESIS	Safety in professional AI is boundary design: clarity about what the system will do, what it will limit and what it will hand to a person, together with evidence that it behaves that way.
THE EVIDENCE GATE	Can the system show how it decides to clarify, limit or escalate, and can your reviewers see and override that behavior?
THE NEXT STEP	Test boundary behavior on your own hardest cases, or request a governance-focused enterprise briefing.
EDITION	Q4 2026. About 16 minutes to read. Free download.

CONTENTS

Read this first.	2
Executive summary	3
The professional-judgment line	4
How AI assistance fails at the boundary	5
Clarify, limit, escalate	5
A worked example: one question, three outcomes	7
Risk classification and decision rights	8
Human oversight that actually works	8
Designing the experience of a boundary	9
Governance, accountability and legal considerations	10
Monitoring and incident response	11
Where ZoikoLogia™ fits	12
Next step: test boundaries on your hardest cases	13
Glossary	14
Important information	14

SUMMARY

Executive summary

Professional accounting work rests on judgment, and judgment is accountable. The central safety question for AI in this setting is not how smart the system is. It is whether it knows where its remit ends.

AI assistance is genuinely useful on the supporting side of professional work: finding authority, comparing positions, drafting and explaining. It becomes unsafe when it quietly crosses into the decision itself, answers definitively where judgment is needed, or leaves a tired reviewer to approve what they cannot properly check.

This paper describes three behaviors a well-bounded system needs, which are to **clarify**, **limit** and **escalate**. It then covers the surrounding governance: how to classify risk, who decides at each level, how to make human oversight real rather than ceremonial, and how to monitor boundaries over time.

Five boundary tests for any AI used in accounting work

#	TEST	THE QUESTION TO PUT TO ANY VENDOR
1	Remit	What is the system for, what is it not for, and where is that written down?
2	Clarify	Does it ask for missing facts and state its assumptions, or answer anyway?
3	Limit	Can it answer the supportable part and say plainly what it cannot conclude?
4	Escalate	Which matters go to a person, with what context, and how does the loop close?
5	Oversight	Can reviewers see the basis for an answer, and can they override it?

THE SHORT VERSION

A system that always answers is not safe in professional work. Safety shows up in what it declines, qualifies and escalates, so test those behaviors first.

What to do with this paper. Use Section 3 to specify the three behaviors, Section 5 to set decision rights, and Section 9 to monitor them once live. If you would like to test these ideas against ZoikoLogia™ with Kriton™ specifically, the last section explains how.

SECTION 1

The professional-judgment line

The line between assistance and judgment is not binary, and treating it as binary is how most boundary failures begin.

Professional judgment in accounting means interpreting facts under applicable requirements, weighing materiality, making estimates and taking positions. It is accountable: a named professional answers for it. AI can support almost every step leading up to it. The risk lies in the last step.

It helps to think of assistance as a spectrum, with the appropriate controls rising as you move toward the decision.

LEVEL OF ASSISTANCE	WHAT THE SYSTEM DOES	TYPICAL CONTROLS
Explain	Describes a concept or requirement in plain language	Source citation; educational framing
Locate	Finds the relevant authority for a question	Approved sources; version and jurisdiction checks
Analyze	Compares positions or applies a requirement to stated facts	Stated assumptions; caveats; reviewer prompt
Recommend	Proposes a treatment or position	Mandatory review; escalation for judgment-heavy matters
Decide	Commits to the outcome	Not a role for the system. Reserved to qualified people

THE BOUNDARY QUESTION

For each workflow, decide in advance the highest level of assistance the system is permitted to reach, and write it down. Boundaries that are implied get crossed.

SECTION 2

How AI assistance fails at the boundary

Boundary failures are rarely dramatic. They are small overreaches that pass review because nothing looked wrong.


FAILURE MODE	WHAT IT LOOKS LIKE	THE CONTROL THAT ADDRESSES IT
Overreach	A definitive answer where judgment is required	Defined remit; limit and escalate behaviors
Underspecified facts	An answer given without the facts it depends on	Clarifying questions; stated assumptions
Jurisdiction drift	Correct text applied to the wrong place or entity	Jurisdiction and applicability checks
Authority confusion	Commentary or internal notes treated as primary authority	Source hierarchy; citations
Automation bias	Reviewers accept output because it looks polished	Evidence-first review design; sampling; calibration
Scope creep	A tool built for one purpose used for another	Permitted-use policy; role-based access

Automation bias deserves special attention. It is a human failure that technology makes more likely, because fluent output invites trust. A system can be technically careful and still unsafe if its design makes uncritical approval the path of least resistance.

SECTION 3

Clarify, limit, escalate

A well-bounded system has three ways of not simply answering. Each is a feature to specify and test.



Clarify

When a question depends on facts the user has not supplied, the system should ask for them or state the assumptions it is making, instead of filling the gap. A confident answer to an under-specified question is the most common source of quiet error.

Limit

Where part of a question can be supported from approved sources and part requires judgment, the system should answer the first and say plainly what it cannot conclude. "Here is what the sources say; whether it applies depends on facts and judgment that need a professional" is a good answer.

Escalate

When a matter is high-risk, unusual or outside the remit, the system hands it to a qualified person with the context they need: the question, the sources consulted and why it was escalated. The original request is not lost, and the outcome is recorded.

Illustrative scenarios

SCENARIO	EXPECTED BEHAVIOR
A user asks how a requirement applies but omits key facts	Clarify: ask for the facts or state assumptions
The facts are clear, but the treatment turns on professional judgment	Limit: set out what the sources say and what needs a professional
The matter is unusual, high-impact or outside approved sources	Escalate: route to a qualified person with context
A user asks the system to simply decide	Decline the decision, explain why, and offer to support the decision-maker

SECTION 4

A worked example: one question, three outcomes

The same starting question can legitimately lead to three different outcomes, depending on the facts and the risk.

Consider a user who asks whether a particular treatment applies to an arrangement. The example is generic and illustrative. What matters is the shape of each outcome.

STARTING POINT	WHAT THE SYSTEM DOES	WHY
Only a one-line description of the arrangement	Clarifies. Asks which type of entity, which period and what the key terms are. If the user wants a provisional view, states its assumptions explicitly	The answer depends on facts that have not been given
Full facts, and a fairly standard pattern	Limits. Sets out what the approved sources say, with citations, notes the points where applying them takes judgment, and prompts a review	Part of the question is supportable and part needs a professional
Full facts, but unusual, material to the client or outside approved sources	Escalates. Routes the matter to a senior reviewer with the question, the sources consulted and the reason, and tells the user it has done so	The risk and novelty exceed what the system should handle alone

In each case the user receives something useful and honest: a question, a qualified answer or a handoff. None of them is a refusal that leaves the user stuck, and none is a confident answer to a question that should not have been answered that way.

What each outcome should record

- The question as asked, and any facts or assumptions supplied.
- The sources consulted and their versions.
- Which route was taken and why.
- For escalations, who received the matter and what they decided.

SECTION 5

Risk classification and decision rights

Behavior should vary with risk, and the people allowed to accept an answer should vary with it too.

Not every question carries the same consequence. Useful factors include the size of the potential impact, how reversible the outcome is, how novel or ambiguous the situation is, how clear the governing authority is, the jurisdiction involved and the sensitivity of the client or matter. The tiers below are illustrative. The principle matters more than the labels.

ILLUSTRATIVE TIER	TYPICAL HANDLING	WHO MAY ACCEPT THE OUTPUT	EVIDENCE EXPECTED
Routine	Answer with citations	The user, within policy	Question, sources, answer
Elevated	Answer with caveats and a review prompt	A reviewer	Plus the caveats shown and reviewer decision
High-risk	Do not answer as routine. Escalate	A qualified senior person	Plus the escalation context and outcome

Write decision rights down. When nobody has specified who may accept which tier, acceptance defaults to whoever is closest to the screen, which is exactly the person least able to assess a high-risk matter.

SECTION 6

Human oversight that actually works

Putting a person in the loop is not the same as having oversight.

Oversight fails quietly when review becomes a click-through. Three design choices make the difference.

- **Evidence first.** Show the sources and the basis for an answer before, not after, the reviewer is asked to approve it.
- **Show uncertainty.** Surface caveats, assumptions and gaps prominently. A reviewer cannot weigh what they cannot see.
- **Require engagement.** For higher tiers, ask the reviewer to record a decision and a reason, not simply to press a button.

Capacity and calibration

Reviewers need time. If escalations arrive faster than people can properly consider them, the process degrades into rubber-stamping. Plan reviewer capacity as part of the design, sample accepted output for quality, and periodically compare reviewers' decisions to detect drift in either direction.

A PRACTICAL TEST

Ask a reviewer to explain, from the screen alone, why an answer is right and which source supports it. If they cannot do that quickly, the review is not doing its job.

SECTION 7

Designing the experience of a boundary

A boundary that users experience as an obstacle will be worked around.

The best boundary logic fails if people dislike meeting it. Users who feel blocked, lectured or left without a next step will route around the system, and the matters that most need a boundary are the ones they will route around it for. Boundary behavior is a user-experience problem as much as a policy problem.

A POOR BOUNDARY EXPERIENCE	A BETTER BOUNDARY EXPERIENCE
A bare refusal with no explanation	A short reason and a clear next step
A caveat buried at the end of a long answer	The limits stated up front, in plain language
An escalation that vanishes into a queue	Confirmation of who has the matter and what happens next
Different wording each time a boundary is hit	Consistent language, so users learn to recognize each behavior

Four practices

- **Explain briefly.** Say why the system is clarifying, limiting or escalating, in a sentence.
- **Always offer a next step.** A boundary that ends the conversation teaches people to stop asking.
- **Make escalation easy.** If handing a matter to a person takes effort, people will skip it.
- **Train with scenarios.** Onboard users with real examples of each behavior, so the first time they meet a boundary is not in a live matter.

Finally, avoid anything that hides limits to make output look more impressive. Polished presentation without visible caveats is a design choice that increases automation bias.

SECTION 8

Governance, accountability and legal considerations

A policy gives the boundaries force. It should be specific enough to apply and short enough to be read.

A workable AI-use policy for accounting typically addresses the following. This is a checklist for your own governance process, not legal advice.

POLICY ELEMENT	WHAT IT SHOULD SETTLE
Permitted uses	Which workflows and which levels of assistance are allowed
Prohibited uses	What the system must not be used for
Source approval	Who approves sources, and how they are kept current
Review requirements	Which tiers require review, by whom, and what is recorded
Data handling	What may be entered, retained and shared, and by whom
Incident handling	How errors and near-misses are reported and learned from
Accountability	Who owns the outcome. Professional responsibility remains with the professional

Questions for legal readers

- How do our professional and contractual obligations apply when AI supports a deliverable?
- What do we need to tell clients, and when? This is a decision for your organization and advisers.
- What records must we keep to demonstrate that required review took place?
- How do vendor terms interact with our confidentiality obligations?

These are questions to take to your own advisers. Nothing in this paper answers them for you.

SECTION 9

Monitoring and incident response

Boundaries drift. What you measure after launch tells you whether they still hold.

WHAT TO MONITOR	WHAT IT CAN TELL YOU
Share of questions clarified, limited or escalated	Whether boundary behavior is occurring at a plausible rate
Reviewer overrides and amendments	Where the system and the professionals disagree
Escalations that reviewers judged unnecessary or late	Whether thresholds are too loose or too tight
Near-misses	Failures that were caught, which are the cheapest lessons
Out-of-remit requests	Whether users are trying to use the tool for other purposes

Treat incidents as a source of design information. For each, ask which boundary failed, which control should have caught it, and whether the answer is a change to the system, the policy or the training. Keep changes under change control, so that tuning a boundary is a recorded decision.

SECTION 10

Where ZoikoLogia™ fits

We built ZoikoLogia™ with Kriton™ around these ideas. You should hold us to the same checks as everyone else.

ZoikoLogia™ with Kriton™ is a governed AI accounting intelligence platform. Its design intent follows the approach in this paper: Kriton™ is designed to clarify, limit or escalate rather than answer definitively, with a governance control architecture that runs from risk classification through to escalation, and four modes (Learning, Workflow, Review and Admin) so that roles see what their work needs.

We would rather you test than take our word. Bring your hardest boundary cases, including ones where the right behavior is to decline. The checks in Sections 3 and 9 apply to us as they do to anyone.

We would rather you verify than trust. Evidence stays with the pages that own it, and each page states what is currently published and what is not:

IF YOU WANT TO CHECK	WHERE TO LOOK
How the organization is governed	Governance
Compliance information and its status	Compliance
How data is protected	Privacy & Security
Provider evaluation and questionnaires	Provider Due Diligence
Accessibility	Accessibility
Documented technical interfaces	API Reference
Recent product changes	Release Notes
Where to start	Trust

WHAT THIS PAPER DOES NOT CLAIM

It does not claim any certification, attestation or compliance status. It makes no statement about pricing, availability or outcomes for your organization, and it provides no accounting, tax, legal, audit or compliance advice. Current evidence and its status are on the Trust pages.

SECTION 11

Next step: test boundaries on your hardest cases

Boundary behavior is easiest to judge on the questions a system should not simply answer.

Bring a mix of routine questions, under-specified ones, judgment-heavy ones and ones that should be escalated, together with the people who would review the results.

Request Enterprise Briefing

The recommended next step. A governance-focused conversation for AI governance, audit, legal and finance stakeholders.

Book a Demo

A guided walkthrough of the product.

Request Pilot

A controlled evaluation of one defined workflow.

Trust

Evidence and specialist pages, with their current status.

Find each on the ZoikoLogia™ website. A submission through any of them does not confirm pricing, availability, response timing or commercial terms.

What to bring to the conversation

BRING	WHY IT HELPS
Your draft AI-use policy, or its gaps	So the boundaries can be compared with it
Ten hard questions across the four behaviors	Routine, clarify, limit and escalate cases
The tiers and decision rights you intend to use	Shows who would accept what
Your reviewers	They judge whether oversight is real
Your incident process	So monitoring can be designed alongside

APPENDIX A

Glossary

TERM	MEANING IN THIS PAPER
Professional judgment	Accountable interpretation of facts under applicable requirements, including estimates and positions.
Remit	The defined purpose and limits of what a system is used for.
Clarify	Asking for missing facts, or stating assumptions, before answering.
Limit	Answering only the supportable part and stating what cannot be concluded.
Escalate	Handing a matter to a qualified person with the context they need.
Risk classification	Sorting matters by consequence so that handling varies accordingly.
Decision rights	Who may accept an output at each risk tier.
Automation bias	The tendency to over-trust automated output, particularly when it is fluent.
Near-miss	An error or boundary failure that was caught before it caused harm.

APPENDIX B

Important information

This paper is educational research. It does not constitute accounting, tax, legal, audit or compliance advice, and professional decisions must remain subject to qualified human review and applicable rules. The tiers, tables, checklists and scorecards are illustrative frameworks to support your own evaluation, not statements about the performance of any product. Statements about ZoikoLogia™ describe its design intent and are not guarantees of outcomes. Verify current status on the relevant pages of the ZoikoLogia™ website.

ZoikoLogia™ and Kriton™ are trademarks of Zoiko Tech Inc. © Zoiko Logia | All rights reserved. Suggested citation: ZoikoLogia™, AI Safety and Professional Boundaries in Accounting, Q4 2026.