

How to Evaluate AI Accounting Outputs Before Deployment

● AI Governance

● CTO

● Procurement

Full paper: 13 min read
Free download

THE QUESTION

How do you know an AI accounting assistant is good enough to deploy?

THE SHORT ANSWER

Evaluate AI accounting output as a decision process, not a score. Define what good means for your work, build a test set that includes the hard cases, measure against a baseline, set go/no-go gates before you see results, and keep evaluating after launch.

FOUR KEY IDEAS

1 Demos and benchmarks are not enough

Each answers a different question. Only your own evaluation shows how a system behaves on your sources, with your reviewers.

2 The test set is the evaluation

Include routine, boundary, no-answer, conflict, jurisdiction and misuse cases. A set without hard cases makes a pass meaningless.

3 Write the gates first

Non-negotiable, quality, evidence, technical and readiness gates, with named deciders. Outcomes are go, conditional go or no-go.

4 Pass is a snapshot

Sources, models and workflows change. Keep a regression set and define triggers for re-evaluation.

WHO SHOULD READ THE FULL PAPER

AI governance owners, CTOs and technical leaders, and procurement teams who must decide whether an AI accounting assistant is good enough to roll out.

INSIDE THE FULL PAPER

What the full paper covers

#	SECTION	WHAT IT ANSWERS
1	Why demos and leaderboards are not evaluation	Each source of evidence answers a different question. None of them answers yours.
2	Define what good means	Accuracy alone is too blunt. Professional work needs several kinds of quality, weighted by risk.
3	Build a test set	The test set is the evaluation. Its composition decides what a pass means.
4	Run the evaluation	A fair run needs a baseline, consistent reviewers and a record of every result.
5	Go/no-go gates	Gates turn evidence into a decision. They only work if they were written first.
6	Procurement and technical due diligence	Run these reviews alongside the evaluation, not after it.
7	Common evaluation mistakes	Most evaluations that mislead do so in one of a few predictable ways.
8	After launch: keep evaluating	A pass is a snapshot. Systems, sources and work all change.

QUESTIONS THIS PAPER HELPS YOU ASK

☐	What does good mean for this workflow, and which dimensions are gates?
☐	Does our test set include boundary and no-answer cases?
☐	Are the go/no-go gates written down, with named decision-makers?
☐	Have procurement and security been involved from the start?
☐	What will trigger re-evaluation after launch?

GO DEEPER

Download the full paper

Free download from the White Paper Library. About 13 minutes to read.

Then bring it into a conversation: **Request Enterprise Briefing**, **Book a Demo** or **Request Pilot** on the ZoikoLogia™ website.

About this summary. It condenses the paper's main ideas and is not a substitute for it. This summary is educational research only; it is not accounting, tax, legal or audit advice, and it claims no certification, attestation or compliance status. A submission through any ZoikoLogia™ call to action does not confirm pricing, availability, response timing or commercial terms.