

WHITE PAPER · IMPLEMENTATION GUIDE

# How to Evaluate AI Accounting Outputs Before Deployment

Evaluation, benchmarks, test cases and go/no-go gates, so buyers can decide with evidence before any wider rollout.

● AI Governance

● CTO

● Procurement

## FORMAT

PDF · Free download

## READING TIME

About 13 minutes

## EDITION

Q4 2026

Educational research only. This paper does not constitute accounting, tax, legal, audit or compliance advice.

ABOUT THIS PAPER

Read this first.

This guide shows buyers how to evaluate AI accounting output before wider rollout. It is practical rather than theoretical: what to define, how to build a test set, how to run the evaluation, how to set gates and what to do after launch. It is vendor-neutral by design, so you can apply it to any product, including ours.

|                   |                                                                                                                                                                                                     |
|-------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| WHO IT IS FOR     | AI governance owners, CTOs and technical leaders, and procurement teams who must decide whether an AI accounting assistant is good enough to roll out.                                              |
| THE QUESTION      | How do you know an AI accounting assistant is good enough to deploy?                                                                                                                                |
| THE THESIS        | Evaluation is a decision process, not a score: define what good means for your work, test it on your own cases against a baseline, set go/no-go gates in advance, and keep evaluating after launch. |
| THE EVIDENCE GATE | Are the go/no-go gates written down before testing begins, with named people who will make the decision?                                                                                            |
| THE NEXT STEP     | Run a controlled evaluation on your own cases, or request an enterprise briefing.                                                                                                                   |
| EDITION           | Q4 2026. About 13 minutes to read. Free download.                                                                                                                                                   |

CONTENTS

|                                                                    |    |
|--------------------------------------------------------------------|----|
| Read this first. . . . .                                           | 2  |
| Executive summary . . . . .                                        | 3  |
| Why demos and leaderboards are not evaluation . . . . .            | 4  |
| Define what good means . . . . .                                   | 5  |
| Build a test set . . . . .                                         | 5  |
| Run the evaluation . . . . .                                       | 6  |
| Go/no-go gates . . . . .                                           | 7  |
| Procurement and technical due diligence . . . . .                  | 8  |
| Common evaluation mistakes . . . . .                               | 9  |
| After launch: keep evaluating . . . . .                            | 9  |
| Where ZoikoLogia™ fits . . . . .                                   | 10 |
| Next step: run a controlled evaluation on your own cases . . . . . | 11 |
| Glossary . . . . .                                                 | 12 |
| Important information . . . . .                                    | 12 |

## SUMMARY

# Executive summary

Most AI evaluations fail in the same way: the criteria are agreed after the results arrive. A good evaluation reverses the order. It decides what good means, what evidence will count and who will decide, and only then looks at the output.

Demos, public benchmarks and reference customers all have a place, but none of them tells you how a system will behave on your work, with your sources, judged by your reviewers. Only your own evaluation does that.

This guide describes a decision process in seven parts, from defining quality, through building a test set and running it, to setting go/no-go gates and monitoring after launch. It also covers what procurement and technical due diligence add, so those reviews run alongside the evaluation instead of after it.

## Five principles of a defensible evaluation

| # | PRINCIPLE                  | WHAT IT MEANS IN PRACTICE                                                  |
|---|----------------------------|----------------------------------------------------------------------------|
| 1 | Define before you test     | Agree what good means, and the gates, before seeing any output             |
| 2 | Use your own cases         | Test on questions and sources from your real work, including the hard ones |
| 3 | Measure against a baseline | Compare to how the work is done today, with the same people                |
| 4 | Test the edges             | Include cases the system should decline, qualify or escalate               |
| 5 | Decide, then keep checking | Record the decision and its reasons, and re-evaluate after change          |

### THE SHORT VERSION

An evaluation is only as good as the questions in it. If the test set contains no hard cases, a pass means very little.

**What to do with this guide.** Use Section 2 to define quality, Section 3 to build a test set, and Section 5 to write your go/no-go gates before you begin. If you want to test these ideas against ZoikoLogia™ with Kriton™ specifically, the last section explains how.

## SECTION 1

# Why demos and leaderboards are not evaluation

Each source of evidence answers a different question. None of them answers yours.

Buyers usually start from what is available: a vendor demonstration, a published benchmark or a reference customer. Each is useful. Each also has a built-in limit.

| EVIDENCE                   | WHAT IT CAN TELL YOU                                  | WHAT IT CANNOT TELL YOU                                   |
|----------------------------|-------------------------------------------------------|-----------------------------------------------------------|
| <b>Vendor demo</b>         | What the product can do at its best                   | How it behaves on questions the vendor did not choose     |
| <b>Public benchmark</b>    | General capability, on a defined task                 | Whether it handles your sources, jurisdictions and risk   |
| <b>Reference customer</b>  | That another organization found it workable           | Whether your work, controls and reviewers are comparable  |
| <b>Your own evaluation</b> | How it behaves on your work, judged by your reviewers | Nothing beyond what you included, so the test set matters |

## A USEFUL REFRAME

Stop asking “how accurate is it?” and start asking “how does it behave on the cases that matter to us, including the ones it should not answer?”.

## SECTION 2

## Define what good means

Accuracy alone is too blunt. Professional work needs several kinds of quality, weighted by risk.

Before testing, write down the dimensions that matter for the workflow you are evaluating, and which of them are non-negotiable. The list below is a starting point.

| DIMENSION                | THE QUESTION IT ANSWERS                                      | HOW YOU MIGHT JUDGE IT                       |
|--------------------------|--------------------------------------------------------------|----------------------------------------------|
| <b>Correctness</b>       | Is the answer right for the facts and the period?            | Reviewer judgment against a reference answer |
| <b>Source support</b>    | Does each claim rest on an approved source that supports it? | Open and check the cited passages            |
| <b>Completeness</b>      | Does it cover what a reviewer would expect?                  | Reviewer checklist per case type             |
| <b>Boundary behavior</b> | Does it decline, qualify or escalate when it should?         | Outcome on boundary and no-answer cases      |
| <b>Consistency</b>       | Do equivalent questions get equivalent answers?              | Paired questions in different wording        |
| <b>Reviewability</b>     | Can a reviewer check it quickly?                             | Time and effort for a reviewer to sign off   |
| <b>Efficiency</b>        | Does it save time on the real workflow?                      | Time against the baseline                    |

Not every dimension has equal weight. A wrong answer in a high-risk matter is not offset by speed in a routine one. State which dimensions are gates, which are scored, and which are simply observed.

## SECTION 3

## Build a test set

The test set is the evaluation. Its composition decides what a pass means.

### Where cases come from

- Real, past questions from your own work, anonymized or approved for use.
- Synthetic cases that mirror real ones without sensitive data.
- Cases written by your subject-matter experts to probe known difficulties.

| CASE CATEGORY           | WHAT IT TESTS                                             |
|-------------------------|-----------------------------------------------------------|
| Routine                 | Whether the basics are reliable                           |
| Boundary                | Whether judgment-heavy matters are qualified or escalated |
| No-answer               | Whether it declines when the sources cannot answer        |
| Conflict                | Whether it surfaces disagreement between approved sources |
| Jurisdiction and period | Whether applicability and effective dates are respected   |
| Misuse                  | Whether it resists being asked to do what it should not   |


## Good practice

- **Write reference answers independently of the vendor.** The people judging should not be the people selling.
- **Start small and grow.** Include each category several times rather than one category many times.
- **Protect data.** Use only data approved for evaluation, and keep the test set under version control.
- **Keep a held-back portion.** A set the vendor has never seen is a better test of generalization.

### SECTION 4

# Run the evaluation

A fair run needs a baseline, consistent reviewers and a record of every result.



Six steps in order. The baseline and the gates are fixed before the system runs.

- **Take the baseline first.** Measure how the work is done today, with the same people and the same kinds of case.
- **Have reviewers apply their normal professional standard.** Do not lower the bar so the pilot looks good.
- **Judge retrieval and answers separately where you can.** A combined score hides which part failed.
- **Record every result by category.** An overall average conceals the failures that matter most.
- **Log escalations and ask whether each was right.** Escalating too much is a cost. Escalating too little is a risk.

Where practical, ask reviewers to assess output without knowing whether it came from the system or from a colleague. This reduces the pull of expectation in either direction.

## SECTION 5

# Go/no-go gates

Gates turn evidence into a decision. They only work if they were written first.

A gate is a condition that must be met before the next step. The categories below are typical. The criteria and thresholds are yours to set, with your risk appetite and advisers.

| GATE TYPE                     | EXAMPLE CRITERION                                                                     | WHO DECIDES                   |
|-------------------------------|---------------------------------------------------------------------------------------|-------------------------------|
| <b>Non-negotiable</b>         | No high-risk test case is answered as routine                                         | AI governance owner           |
| <b>Quality</b>                | Answers meet the standard you set for correctness and source support in each category | Accountable professional lead |
| <b>Evidence</b>               | An auditor can reconstruct a sample of answers from the record                        | Audit or assurance lead       |
| <b>Technical and security</b> | Security and data-handling review completed and accepted                              | CTO and security lead         |
| <b>Readiness</b>              | Reviewers have capacity and training for the expected volume                          | Workflow owner                |

## Three possible outcomes

- **Go.** All gates are met, and the decision and reasons are recorded.
- **Conditional go.** Proceed with named controls, such as a narrower scope or mandatory review, and a date to re-assess.
- **No-go.** A gate is not met. Record why, and what would change the answer.

## SECTION 6

# Procurement and technical due diligence

Run these reviews alongside the evaluation, not after it.

Strong quality results can still end at procurement or security review if those teams were brought in late. Involve them at the start, with the questions they will ask.

| FUNCTION                    | WHAT THEY WILL WANT TO KNOW                                                        |
|-----------------------------|------------------------------------------------------------------------------------|
| <b>Procurement</b>          | Vendor information, contracting process, data handling terms, portability and exit |
| <b>Security and IT</b>      | Architecture, access control, data handling, and where to find current evidence    |
| <b>AI governance</b>        | Boundary behavior, escalation, monitoring and policy fit                           |
| <b>Audit</b>                | What record is produced, and whether it supports reconstruction                    |
| <b>Legal and compliance</b> | Obligations, confidentiality and the vendor's contractual position                 |

For vendor evidence, use the vendor's own due-diligence and trust pages, which should state what is currently published and what is not.

## SECTION 7

# Common evaluation mistakes

Most evaluations that mislead do so in one of a few predictable ways.

These mistakes are easy to make because each feels reasonable at the time. Naming them in advance is the cheapest protection.

| MISTAKE                                       | WHAT IT CAUSES                                      | THE FIX                                                      |
|-----------------------------------------------|-----------------------------------------------------|--------------------------------------------------------------|
| <b>Choosing criteria after seeing results</b> | Criteria drift toward whatever the system does well | Write quality dimensions and gates first                     |
| <b>Testing only routine cases</b>             | A pass that hides weakness at the edges             | Include boundary, no-answer and conflict cases               |
| <b>Letting the vendor pick the questions</b>  | Results that flatter the product                    | Use your own cases, plus a held-back set                     |
| <b>Reporting one overall score</b>            | Serious failures averaged away                      | Report by category, and treat non-negotiables as gates       |
| <b>Using unqualified reviewers</b>            | Errors missed, or valid answers rejected            | Have qualified professionals judge, to their normal standard |
| <b>Skippping the baseline</b>                 | No honest way to judge improvement                  | Measure today's process first                                |
| <b>Treating a pass as permanent</b>           | Unnoticed drift after changes                       | Keep a regression set and re-evaluation triggers             |

A useful discipline is to ask, before every evaluation, "what result would make us stop?". If nobody can answer, the gates have not really been written.

## SECTION 8

# After launch: keep evaluating

A pass is a snapshot. Systems, sources and work all change.

| TRIGGER FOR RE-EVALUATION                      | WHY                                                        |
|------------------------------------------------|------------------------------------------------------------|
| <b>Approved sources change</b>                 | Answers may change, and earlier results may no longer hold |
| <b>The model or system is updated</b>          | Behavior can shift in ways release notes do not capture    |
| <b>The workflow or scope expands</b>           | New cases bring new risks                                  |
| <b>A significant error or near-miss occurs</b> | It reveals a gap the test set did not cover                |
| <b>A scheduled review date arrives</b>         | Evidence ages even when nothing visible changes            |

Keep a regression set of earlier cases and re-run it after any significant change. Record the outcome, and treat changes to boundaries, sources or models as decisions under change control, not as quiet updates.

## SECTION 9

# Where ZoikoLogia™ fits

We built ZoikoLogia™ with Kriton™ around these ideas. You should hold us to the same checks as everyone else.

ZoikoLogia™ with Kriton™ is a governed AI accounting intelligence platform, designed around source-backed answers, risk-aware handling and escalation, and qualified people who stay accountable. Those are the behaviors an evaluation should probe.

We would rather you test than take our word. Bring your hardest cases, including ones where the right behavior is to decline. The approach in this guide applies to us as it does to anyone.

We would rather you verify than trust. Evidence stays with the pages that own it, and each page states what is currently published and what is not:

| IF YOU WANT TO CHECK                   | WHERE TO LOOK          |
|----------------------------------------|------------------------|
| How the organization is governed       | Governance             |
| Compliance information and its status  | Compliance             |
| How data is protected                  | Privacy & Security     |
| Provider evaluation and questionnaires | Provider Due Diligence |
| Accessibility                          | Accessibility          |
| Documented technical interfaces        | API Reference          |
| Recent product changes                 | Release Notes          |
| Where to start                         | Trust                  |

### WHAT THIS PAPER DOES NOT CLAIM

It does not claim any certification, attestation or compliance status. It makes no statement about pricing, availability or outcomes for your organization, and it provides no accounting, tax, legal, audit or compliance advice. Current evidence and its status are on the Trust pages.

## SECTION 10

# Next step: run a controlled evaluation on your own cases

The most informative evidence is how a system behaves on your work, judged by your reviewers.

Bring a draft test set, the dimensions that matter, and your draft gates. A short, structured evaluation will tell you more than any amount of demonstration.

**Request Pilot**

The recommended next step. A controlled evaluation of one defined workflow.

**Request Enterprise Briefing**

Decision context for governance, technical, procurement and finance stakeholders.

**Book a Demo**

A guided walkthrough of the product.

**Provider Due Diligence**

Provider evidence and questionnaire support.

Find each on the ZoikoLogia™ website. A submission through any of them does not confirm pricing, availability, response timing or commercial terms.

## What to bring to the conversation

| BRING                                                  | WHY IT HELPS                                      |
|--------------------------------------------------------|---------------------------------------------------|
| <b>A draft test set across the six case categories</b> | The evaluation is only as good as its cases       |
| <b>Your quality dimensions, with gates marked</b>      | So what counts as a pass is agreed first          |
| <b>A baseline for the workflow</b>                     | Time and quality before, for an honest comparison |
| <b>Your reviewers and decision-makers</b>              | They judge, and they decide                       |
| <b>Your procurement and security questions</b>         | So reviews run in parallel                        |

## APPENDIX A

# Glossary

| TERM                    | MEANING IN THIS PAPER                                                              |
|-------------------------|------------------------------------------------------------------------------------|
| <b>Baseline</b>         | How the work is done today, measured before the evaluation, for comparison.        |
| <b>Test set</b>         | The collection of cases used to evaluate a system.                                 |
| <b>Reference answer</b> | The answer a qualified person judges correct for a test case.                      |
| <b>Go/no-go gate</b>    | A condition that must be met before proceeding to the next step.                   |
| <b>Boundary case</b>    | A case where judgment is needed and the right behavior is to qualify or escalate.  |
| <b>No-answer case</b>   | A case the approved sources cannot answer, where the right behavior is to decline. |
| <b>Regression set</b>   | Earlier cases re-run after a change to check results have not worsened.            |
| <b>Held-back set</b>    | Cases kept out of tuning or demonstration so they test generalization.             |
| <b>Conditional go</b>   | A decision to proceed subject to named controls and a re-assessment date.          |

## APPENDIX B

# Important information

This paper is educational research. It does not constitute accounting, tax, legal, audit or compliance advice, and professional decisions must remain subject to qualified human review and applicable rules. The tiers, tables, checklists and scorecards are illustrative frameworks to support your own evaluation, not statements about the performance of any product. Statements about ZoikoLogia™ describe its design intent and are not guarantees of outcomes. Verify current status on the relevant pages of the ZoikoLogia™ website.

ZoikoLogia™ and Kriton™ are trademarks of Zoiko Tech Inc. © Zoiko Logia | All rights reserved. Suggested citation: ZoikoLogia™, How to Evaluate AI Accounting Outputs Before Deployment, Q4 2026.