

WHITE PAPER · TECHNICAL PAPER

Source-Backed RAG for Accounting Workflows

Why accounting AI must cite, rank and limit sources, and how retrieval-augmented generation should be governed rather than left to model memory alone.

CTO

Accounting Firm

Enterprise Finance

FORMAT

PDF · Free download

READING TIME

About 14 minutes

EDITION

Q4 2026

Educational research only. This paper does not constitute accounting, tax, legal, audit or compliance advice.

ABOUT THIS PAPER

Read this first.

This paper explains how retrieval-augmented generation works in accounting settings, where it fails, and what governing it involves. It is written for technical and finance leaders who must decide whether a retrieval-based assistant is fit for professional work. It is vendor-neutral by design, so you can use it to evaluate any approach, including ours.

WHO IT IS FOR	CTOs and technical leaders, accounting firm leaders and enterprise finance teams, and the engineers and reviewers who will evaluate retrieval-based AI.
THE QUESTION	Should an accounting AI answer from model memory, or from sources it can cite?
THE THESIS	Retrieval-augmented generation (RAG) helps accounting AI only when retrieval itself is governed: which sources are eligible, how they are ranked, how much reaches the model, and what happens when nothing relevant is found.
THE EVIDENCE GATE	For any answer, can the system show which approved source passages were retrieved and why, and does it stop when none are good enough?
THE NEXT STEP	Test retrieval on your own questions and sources, or request an enterprise briefing.
EDITION	Q4 2026. About 14 minutes to read. Free download.

CONTENTS

Read this first.	2
Executive summary	3
Why accounting cannot run on model memory	4
How ungoverned RAG fails in accounting	5
Governing the corpus	6
Governing retrieval and ranking	7
Governing generation and citation	7
Evaluating a source-backed system	8
Architecture questions for technical leaders	9
Where ZoikoLogia™ fits	10
Next step: test retrieval on your own questions	11
Glossary	12
Important information	12

SUMMARY

Executive summary

Retrieval-augmented generation is the most common way to make an AI assistant answer from documents instead of from memory. In accounting, the interesting question is not whether to use it. It is whether the retrieval is governed.

RAG works in two steps. The system first retrieves passages from a body of source material, then gives those passages to a language model to compose an answer. Done well, an answer is grounded in something a reviewer can open. Done badly, it is a fluent answer attached to the wrong passage, an out-of-date version or a different jurisdiction.

This paper argues that the quality of a RAG system is decided less by the model than by the rules around it: what may be retrieved, how results are ranked, how much context is passed on, how answers cite their support, and what the system does when retrieval comes back empty.

Five properties of governed retrieval

#	PROPERTY	THE QUESTION TO PUT TO ANY VENDOR
1	Eligible sources	Which sources may be retrieved, who approved them, and can that list be seen and changed by the right people?
2	Authority-aware ranking	Does ranking reflect authority, applicability and effective date, or only textual similarity?
3	Limited context	How many passages reach the model, and what stops irrelevant ones from diluting the answer?
4	Verified citation	Does each substantive claim point to a passage that actually supports it?
5	Honest abstention	When nothing relevant is found, does the system say so and escalate instead of improvising?

THE SHORT VERSION

RAG is a pattern, not a guarantee. Ask to see retrieval behave on the questions it should not be able to answer. That is where governance shows.

What to do with this paper. Use Section 7 as a question list for your technical evaluation, and Section 6 to design a test that separates retrieval quality from answer quality. If you want to test these ideas against ZoikoLogia™ with Kriton™ specifically, the last section explains how.

SECTION 1

Why accounting cannot run on model memory


A language model's memory is a poor place to keep rules that change, vary by jurisdiction and carry consequences.

A language model learns patterns from a very large body of text during training. What it "knows" is frozen at a point in time, blended from many sources, and not attributed to any of them. For general questions that is often fine. For accounting it creates four problems.

- **Currency.** Requirements change, and an answer that was right for one period may be wrong for the next.
- **Applicability.** The right answer depends on jurisdiction, entity type and framework, none of which memory tracks reliably.
- **Provenance.** A memorised answer has no source a reviewer can open, so it cannot be checked, only trusted.
- **Calibration.** A model sounds equally confident whether or not it is on firm ground.

What retrieval-augmented generation is

RAG changes the question the model is asked. Instead of "what do you remember about this?", the system first finds relevant passages in an approved body of material, then asks the model to answer using only those passages and to point back to them.



Five steps, each of which can be governed. Most quality problems originate in the second and third.

The pattern is powerful, but it is not a guarantee. Retrieval can fetch the wrong passage. The model can misread or ignore a good one. The underlying collection can be stale. Governance exists to make those failures visible and rare.

A USEFUL REFRAME

Stop asking "does it use RAG?" and start asking "what governs the retrieval?". Nearly every serious product can say yes to the first.

SECTION 2

How ungoverned RAG fails in accounting

Most failures are not the model inventing things. They are retrieval returning the wrong thing, convincingly.

The table lists failure modes that matter in professional accounting work, and the governance response to each. None of them requires a bad model. They arise from how retrieval is set up.

FAILURE MODE	WHAT IT LOOKS LIKE	GOVERNANCE RESPONSE
Wrong-source retrieval	A passage that mentions the right words but governs a different matter	Authority-aware ranking and applicability filters
Stale version	An answer drawn from a superseded document	Versioning and effective-date metadata on every source
Jurisdiction mismatch	Correct text applied in the wrong jurisdiction	Jurisdiction tagging and filtering before ranking
Broken context	A rule retrieved without its exception or definition	Chunking that keeps rules, exceptions and definitions together
Diluted context	Too many marginal passages crowd out the one that matters	Fewer, better passages, with relevance thresholds
Citation theatre	A citation that does not actually support the claim	Claim-level citation and a check that the passage supports it
Silent gap	An answer produced when nothing relevant was found	Explicit abstention, with escalation to a person

Two of these deserve emphasis. **Citation theatre** is dangerous because a footnote lends credibility whether or not the source supports the statement. And the **silent gap** is the most common route from a good demo to a bad outcome, because demos are built on questions the sources can answer.

SECTION 3

Governing the corpus

Retrieval can only be as good as what it is allowed to retrieve.

Source eligibility

The first control is the simplest and the most often skipped: a defined list of sources that may be retrieved. Which sources are approved is a decision for the organization and its professional advisers. A governed system enforces that decision, and makes it visible and changeable by the right people.

Source bundles

In practice, sources are grouped into curated sets for a workflow, a jurisdiction or a client type. Bundling keeps retrieval within a relevant, approved perimeter and makes changes easier to review, because you change a bundle rather than a thousand documents.

Metadata that retrieval depends on

METADATA	WHY RETRIEVAL NEEDS IT
Authority level	Lets primary authority outrank commentary and internal notes
Version and effective date	Lets an answer reflect the right version for the period in question
Jurisdiction and applicability	Prevents correct text being applied in the wrong place or to the wrong entity
Source identity	Makes a citation openable and an audit trail possible
Access permissions	Ensures people only retrieve what they are entitled to see

SOURCE-BACKED IS NOT SOURCE-LISTED

A vendor that lists the sources it “supports” is telling you what was ingested. It is not telling you whether answers are bound to those sources. Test the behavior, not the list.

SECTION 4

Governing retrieval and ranking

The ranking step decides which passages the model ever sees, so it decides most of what the answer can be.

Retrieval systems commonly combine keyword matching, which is good at exact terms and identifiers, with semantic matching, which is good at meaning. Accounting benefits from both, because precise defined terms and loosely worded questions meet in the same query.

Four ranking principles for accounting

- **Authority before similarity.** A passage from primary authority should outrank a closer-worded passage from commentary.
- **Filter before you rank.** Remove passages that do not apply by jurisdiction, entity or period, instead of hoping they rank low.
- **Keep meaning intact.** Split documents so a rule travels with its exception and its definition.
- **Fewer, better passages.** More context is not better context. Passing marginal text to the model invites it to use it.

Knowing when to stop

A governed system sets a relevance threshold. When the best retrieved passages fall below it, the right behavior is to say that nothing adequate was found and, where the matter is risky, to escalate. This is the single behavior most worth testing, because it is the one demos avoid.

SECTION 5

Governing generation and citation

Once passages are chosen, the remaining question is whether the answer stays within them.

- **Answer from the passages.** The model is instructed to rely on the retrieved material and to say what it cannot conclude from it.
- **Cite at claim level.** A single footnote at the end of a paragraph hides which statement rests on which source.
- **Verify the support.** A check confirms each cited passage actually supports the claim attached to it.
- **Quote rules, paraphrase carefully.** Where exact wording matters, quote it. Where it does not, make sure a paraphrase has not changed the meaning.
- **Surface conflict.** If approved sources disagree, show the disagreement. Do not quietly choose one.

None of these requires a more capable model. They are rules around the model, and they are the difference between a citation that decorates an answer and one that holds it up.

THE REVIEWER'S TEST

Pick five answers at random. For each, open the cited passage and ask: does it say what the answer claims? If you cannot do this quickly, reviewers will not do it at all.

SECTION 6

Evaluating a source-backed system

Retrieval quality and answer quality are different things. Measure them separately.

A system can retrieve the right passage and still answer badly, or retrieve the wrong one and still sound right. Combining the two into a single score hides which part failed, and therefore what to fix.

WHAT TO TEST	HOW	WHY TEST IT SEPARATELY
Retrieval	For each test question, did the right passage appear among those retrieved?	Shows whether the corpus, metadata and ranking are working
Answer quality	Given good passages, was the answer correct and complete?	Shows whether generation rules are working
Citation support	Does each cited passage support its claim?	Catches citation theatre
Abstention	Does the system decline or escalate when it should?	Catches the silent gap
Regression	After a source update, do earlier answers change as expected?	Shows whether updates are controlled

Build a test set that includes the edges

- Questions with known answers in approved sources.
- Questions the sources **cannot** answer, which should produce abstention.
- Questions where two sources conflict, or where the right answer depends on jurisdiction or period.
- Questions that should be escalated to a person.

Agree the success criteria before you run the test, and have reviewers apply their normal professional standard.

SECTION 7

Architecture questions for technical leaders

Ten questions that separate a governed deployment from a demonstration.

QUESTION	WHY IT MATTERS
Who approves sources, and can we see and change the list?	Source eligibility is the first control
How are versions and effective dates handled?	Currency decides whether answers are right for the period
How is jurisdiction represented and filtered?	Correct text in the wrong place is still wrong
How are documents split, and do rules keep their exceptions?	Broken context is a common hidden failure
How is ranking decided, and does it weigh authority?	Similarity alone is not authority
What happens when nothing relevant is found?	The most important behavior to observe
Can each claim be traced to a passage, and is support checked?	Distinguishes real citation from decoration
Are retrieved passages and versions logged for each answer?	Without it, an answer cannot be reconstructed later
How are corpora separated between organizations, and who can see what?	A security and confidentiality question
How do source updates roll out, and how are regressions caught?	Change control for knowledge, not only for code

For security, privacy and technical-interface questions, use the owning pages rather than a vendor summary. Section 8 lists them.

SECTION 8

Where ZoikoLogia™ fits

We built ZoikoLogia™ with Kriton™ around these ideas. You should hold us to the same checks as everyone else.

ZoikoLogia™ with Kriton™ is a governed AI accounting intelligence platform. In the terms of this paper, it is designed around governed retrieval: answers drawn from approved sources, with versioning and citation discipline, rather than from model memory alone.

How source bundles, ranking, citation and abstention behave in your deployment is something to verify on your own questions, not to assume from a description. The checks in Sections 6 and 7 apply to us as they do to anyone.

We would rather you verify than trust. Evidence stays with the pages that own it, and each page states what is currently published and what is not:

IF YOU WANT TO CHECK	WHERE TO LOOK
How the organization is governed	Governance
Compliance information and its status	Compliance
How data is protected	Privacy & Security
Provider evaluation and questionnaires	Provider Due Diligence
Accessibility	Accessibility
Documented technical interfaces	API Reference
Recent product changes	Release Notes
Where to start	Trust

WHAT THIS PAPER DOES NOT CLAIM

It does not claim any certification, attestation or compliance status. It makes no statement about pricing, availability or outcomes for your organization, and it provides no accounting, tax, legal, audit or compliance advice. Current evidence and its status are on the Trust pages.

SECTION 9

Next step: test retrieval on your own questions

The most informative evidence is how retrieval behaves on your sources, with your reviewers.

Bring a short list of real questions, including a few the sources should not be able to answer. Watching what the system does with those tells you more than any description of how it works.

Request Enterprise Briefing

The recommended next step. Decision context for engineering, finance, security and governance stakeholders.

Book a Demo

A guided walkthrough of the product.

Request Pilot

A controlled evaluation of one defined workflow.

API Reference

Documented technical interfaces, where published.

Find each on the ZoikoLogia™ website. A submission through any of them does not confirm pricing, availability, response timing or commercial terms.

What to bring to the conversation

BRING	WHY IT HELPS
A list of 15 to 20 real questions	Include routine, boundary and no-answer cases
The sources you would want approved	Source eligibility is the first control
One workflow to measure	Keeps the evaluation narrow enough to prove something
Your jurisdictions and periods	So applicability can be tested
The engineers and reviewers who will judge it	They decide whether retrieval and answers are good enough

APPENDIX A

Glossary

TERM	MEANING IN THIS PAPER
Retrieval-augmented generation (RAG)	Retrieving relevant passages from a body of source material and giving them to a language model so its answer is grounded in them.
Corpus	The body of source material a system may retrieve from.
Source bundle	A curated, approved set of sources grouped for a workflow, jurisdiction or client type.
Chunk	A passage of a document that is stored and retrieved as a unit.
Ranking	Ordering retrieved passages by relevance and, in a governed system, authority and applicability.
Grounding	Constraining an answer to the retrieved material.
Citation	A pointer from a statement to the specific passage that supports it.
Abstention	Declining to answer, or escalating, when nothing adequate was retrieved.
Effective date	The date from which a source applies, used to select the right version for a period.

APPENDIX B

Important information

This paper is educational research. It does not constitute accounting, tax, legal, audit or compliance advice, and professional decisions must remain subject to qualified human review and applicable rules. The tiers, tables, checklists and scorecards are illustrative frameworks to support your own evaluation, not statements about the performance of any product. Statements about ZoikoLogia™ describe its design intent and are not guarantees of outcomes. Verify current status on the relevant pages of the ZoikoLogia™ website.

ZoikoLogia™ and Kriton™ are trademarks of Zoiko Tech Inc. © Zoiko Logia | All rights reserved. Suggested citation: ZoikoLogia™, Source-Backed RAG for Accounting Workflows, Q4 2026.